

```

import tensorflow as tf
from tensorflow.keras.preprocessing.image import ImageDataGenerator from
tensorflow.keras.models import Sequential
from tensorflow.keras.layers import Conv2D, MaxPooling2D, Flatten, Dense, Dropout
from tensorflow.keras.optimizers import Adam train_datagen =
ImageDataGenerator(
    rescale=1./255,      # Normalize pixel values to [0,1]
    rotation_range=20,   # Randomly rotate images by up to 20
degrees width_shift_range=0.2, # Randomly shift images horizontally
height_shift_range=0.2, # Randomly shift images vertically
shear_range=0.2, # Shear transformation zoom_range=0.2, # Zoom in
randomly horizontal_flip=True, # Flip images horizontally
fill_mode='nearest' # Fill pixels after transformations
)
validation_datagen = ImageDataGenerator(rescale=1./255) train_generator =
train_datagen.flow_from_directory(
r'C:\Users\MGM\Desktop\Shweta\Animal_data\Animal_data\train' , target_size=(150,
150), batch_size=32, class_mode='binary' # since only 1 class
)
validation_generator = validation_datagen.flow_from_directory(
r'C:\Users\MGM\Desktop\Shweta\Animal_data\Animal_data\validation'
,
target_size=(150, 150), batch_size=32, class_mode='binary' # since only 1
class
)
Found 123 images belonging to 1 classes. Found 2500
images belonging to 1 classes.
model = Sequential([
    Conv2D(32, (3, 3), activation='relu', input_shape=(150, 150, 3)),
# 32 filters, 3x3 kernel, ReLU activation
    MaxPooling2D(2, 2), # Downsample by factor of 2
    Conv2D(64, (3, 3), activation='relu'), MaxPooling2D(2, 2),
    Conv2D(128, (3, 3), activation='relu'),
    MaxPooling2D(2, 2),
    Flatten(),      # Flatten 3D feature maps to 1D feature vector
    Dropout(0.5),    # Dropout to prevent overfitting

```

```

        Dense(512, activation='relu'), # Fully connected layer with 512 neurons
        Dense(train_generator.num_classes, activation='softmax') # Output layer,
        number of neurons = number of classes ])
C:\Users\MGM\anaconda3\Lib\site-packages\keras\src\layers\
convolutional\base_conv.py:113: UserWarning: Do not pass an
`input_shape` / `input_dim` argument to a layer. When using Sequential models,
prefer using an `Input(shape)` object as the first layer in the model instead.
super().__init__(activity_regularizer=activity_regularizer,
**kwargs) model.compile(loss='categorical_crossentropy', # Suitable loss
for multi-class
classification optimizer=Adam(learning_rate=0.001), # Adam
optimizer with
learning rate 0.001 metrics=['accuracy'] # Track
accuracy metric
) history = model.fit(
train_generator,
    steps_per_epoch=train_generator.samples //
train_generator.batch_size, # Number of batches per epoch epochs=20, #
    Number of training epochs validation_data=validation_generator,
    validation_steps=validation_generator.samples //
validation_generator.batch_size
)
C:\Users\MGM\anaconda3\Lib\site-packages\keras\src\trainers\
data_adapters\py_dataset_adapter.py:121: UserWarning: Your `PyDataset` class
should call `super().__init__(**kwargs)` in its constructor.
`**kwargs` can include `workers`, `use_multiprocessing`,
`max_queue_size`. Do not pass these arguments to `fit()`, as they will be ignored.
self._warn_if_super_not_called()
Epoch 1/20
C:\Users\MGM\anaconda3\Lib\site-packages\keras\src\ops\nn.py:944:
UserWarning: You are using a softmax over axis -1 of a tensor of shape (None, 1). This
axis has size 1. The softmax operation will always return the value 1, which is likely
not what you intended. Did you mean to use a sigmoid instead?
warnings.warn(
C:\Users\MGM\anaconda3\Lib\site-packages\keras\src\losses\losses.py:33:
SyntaxWarning: In loss categorical_crossentropy, expected y_pred.shape to be
(batch_size, num_classes) with num_classes > 1. Received: y_pred.shape=(None, 1).
Consider using 'binary_crossentropy' if you only have 2 classes. return self.fn(y_true,
y_pred, **self.fn_kwargs)

```

3/3 ————— 7s 3s/step – accuracy: 0.0000e+00 –
loss:

0.0000e+00 – val_accuracy: 0.0000e+00 – val_loss: 0.0000e+00

Epoch 2/20

1/3 ————— 0s 148ms/step – accuracy: 0.0000e+00
– loss:

0.0000e+00

C:\Users\MGM\anaconda3\Lib\site-packages\keras\src\trainers\
epoch_iterator.py:116: UserWarning: Your input ran out of data; interrupting training.
Make sure that your dataset or generator can generate at least `steps_per_epoch *
epochs` batches. You may need to use the `.repeat()` function when building your
dataset. self._interrupted_warning()

3/3 ————— 4s 2s/step – accuracy: 0.0000e+00 –
loss:

0.0000e+00 – val_accuracy: 0.0000e+00 – val_loss: 0.0000e+00

Epoch 3/20

3/3 ————— 4s 2s/step – accuracy: 0.0000e+00 –
loss:

0.0000e+00 – val_accuracy: 0.0000e+00 – val_loss: 0.0000e+00

Epoch 4/20

3/3 ————— 4s 2s/step – accuracy: 0.0000e+00 –
loss:

0.0000e+00 – val_accuracy: 0.0000e+00 – val_loss: 0.0000e+00

Epoch 5/20

3/3 ————— 4s 2s/step – accuracy: 0.0000e+00 –
loss:

0.0000e+00 – val_accuracy: 0.0000e+00 – val_loss: 0.0000e+00

Epoch 6/20

3/3 ————— 4s 2s/step – accuracy: 0.0000e+00 –
loss:

0.0000e+00 – val_accuracy: 0.0000e+00 – val_loss: 0.0000e+00

Epoch 7/20

3/3 ————— 4s 2s/step – accuracy: 0.0000e+00 –
loss:

0.0000e+00 – val_accuracy: 0.0000e+00 – val_loss: 0.0000e+00

Epoch 8/20

3/3 ————— 4s 2s/step - accuracy: 0.0000e+00 -
loss:

0.0000e+00 - val_accuracy: 0.0000e+00 - val_loss: 0.0000e+00

Epoch 9/20

3/3 ————— 5s 2s/step - accuracy: 0.0000e+00 -
loss:

0.0000e+00 - val_accuracy: 0.0000e+00 - val_loss: 0.0000e+00

Epoch 10/20

3/3 ————— 4s 2s/step - accuracy: 0.0000e+00 -
loss:

0.0000e+00 - val_accuracy: 0.0000e+00 - val_loss: 0.0000e+00

Epoch 11/20

3/3 ————— 5s 2s/step - accuracy: 0.0000e+00 -
loss:

0.0000e+00 - val_accuracy: 0.0000e+00 - val_loss: 0.0000e+00

Epoch 12/20

3/3 ————— 4s 2s/step - accuracy: 0.0000e+00 -
loss:

0.0000e+00 - val_accuracy: 0.0000e+00 - val_loss: 0.0000e+00

Epoch 13/20

3/3 ————— 5s 2s/step - accuracy: 0.0000e+00 -
loss:

0.0000e+00 - val_accuracy: 0.0000e+00 - val_loss: 0.0000e+00

Epoch 14/20

3/3 ————— 4s 2s/step - accuracy: 0.0000e+00 -
loss:

0.0000e+00 - val_accuracy: 0.0000e+00 - val_loss: 0.0000e+00

Epoch 15/20

3/3 ————— 5s 2s/step - accuracy: 0.0000e+00 -
loss:

0.0000e+00 - val_accuracy: 0.0000e+00 - val_loss: 0.0000e+00

Epoch 16/20

3/3 ————— 4s 2s/step - accuracy: 0.0000e+00 -
loss:

0.0000e+00 - val_accuracy: 0.0000e+00 - val_loss: 0.0000e+00

Epoch 17/20

3/3 ————— 5s 2s/step - accuracy: 0.0000e+00 -
loss:

0.0000e+00 - val_accuracy: 0.0000e+00 - val_loss: 0.0000e+00

Epoch 18/20

3/3 ————— 4s 2s/step - accuracy: 0.0000e+00 -
loss:

0.0000e+00 - val_accuracy: 0.0000e+00 - val_loss: 0.0000e+00

Epoch 19/20

3/3 ————— 5s 2s/step - accuracy: 0.0000e+00 -
loss:

0.0000e+00 - val_accuracy: 0.0000e+00 - val_loss: 0.0000e+00

Epoch 20/20

3/3 ————— 4s 2s/step - accuracy: 0.0000e+00

- loss: 0.0000e+00 - val_accuracy: 0.0000e+00 - val_loss: 0.0000e+00

model.save('animal_classifier_model.h5')

WARNING:absl:You are saving your model as an HDF5 file via
`model.save()` or `keras.saving.save_model(model)`. This file format is considered
legacy. We recommend using instead the native Keras format, e.g.
`model.save('my_model.keras')` or `keras.saving.save_model(model,
'my_model.keras')`.

```
import numpy as np from tensorflow.keras.preprocessing
```

```
import image def predict_animal(img_path):
```

```
    # Load image with target size img =  
    image.load_img(img_path, target_size=(150, 150))  
    # Convert image to array img_array  
    = image.img_to_array(img) # Scale  
    pixel values to [0,1] img_array =  
    img_array / 255.0  
    # Expand dims to add batch size dimension  
    img_array = np.expand_dims(img_array, axis=0)  
  
    # Predict class probabilities predictions  
    = model.predict(img_array)  
    # Get the index of the highest probability  
    predicted_class_index = np.argmax(predictions)
```

```

# Map class index to class label
class_labels      =      list(train_generator.class_indices.keys())
predicted_label   =   class_labels[predicted_class_index]   return
predicted_label

img_path = r"C:\Users\MGM\Desktop\Shweta\Animal_data\Animal_data\
validation\cats\cat.119.jpg"
print(f"Predicted animal: {predict_animal(img_path)}")
1 / 1

```

```

0s 73ms/step
Predicted animal: dolphin
C:\Users\MGM\anaconda3\Lib\site-packages\keras\src\ops\nn.py:944:
UserWarning: You are using a softmax over axis -1 of a tensor of shape (1, 1). This
axis has size 1. The softmax operation will always return the value 1, which is likely
not what you intended. Did you mean to use a sigmoid instead? warnings.warn(
img_path = rimg_path = r"C:\Users\MGM\Desktop\Shweta\Animal_data\
Animal_data\validation\Cats\Cats.119.jpg"
print(f"Predicted animal: {predict_animal(img_path)}") print(f"Predicted animal:
{predict_animal(img_path)}")
1 / 1

```

```

0s 36ms/step
Predicted animal: dolphin
1 / 1

```

```

0s 34ms/step
Predicted animal: dolphin
import matplotlib.pyplot as plt
from tensorflow.keras.preprocessing import image
import numpy as np def predict_and_show(img_path):
    # Load and preprocess image
    img = image.load_img(img_path, target_size=(150, 150))
    img_array = image.img_to_array(img) / 255.0 img_array =
    np.expand_dims(img_array, axis=0)
    # Predict
    predictions = model.predict(img_array) predicted_class_index =
    np.argmax(predictions) class_labels =

```

```
list(train_generator.class_indices.keys()) predicted_label =  
class_labels[predicted_class_index]
```

```
# Show image and prediction plt.imshow(img)  
plt.title(f'Predicted: {predicted_label}') plt.axis('off')  
plt.show()
```

```
# Example usage
```

```
img_path = r"C:\Users\MGM\Desktop\Shweta\Animal_data\Animal_data\validation\Cats\Cat.111.jpg" predict_and_show(img_path)
```

1 / 1 _____ 0s 28ms/step

Predicted : cats



```
from sklearn.metrics import
confusion_matrix, accuracy_score, classification_report
print(accuracy_score) print(confusion_matrix)

<function accuracy_score at 0x0000020D34A49EE0>
<function confusion_matrix at 0x0000020D34A4A020>
<function classification_report at 0x0000020D34A4B380>
print(classification_report)

from sklearn.metrics import confusion_matrix, accuracy_score,
classification_report import numpy as np

# 1. Get true labels from validation generator y_true =
validation_generator.classes # true class indices

# 2. Predict on validation data
```



```

y_pred_probs = model.predict(validation_generator)

# 3. Convert predicted probabilities to class indices if
validation_generator.class_mode == 'binary':
    y_pred = np.argmax(y_pred_probs, axis=1)
else: # for 'binary' classification y_pred = (y_pred_probs >
    0.5).astype(int).reshape(-1)

# 4. Print metrics
print("Accuracy:", accuracy_score(y_true, y_pred))
print("Confusion Matrix:\n", confusion_matrix(y_true, y_pred)) print("Classification
Report:\n", classification_report(y_true, y_pred,
target_names=list(validation_generator.class_indices.keys())))

```

79/79 —————

4s 49ms/step

Accuracy: 1.0

Confusion Matrix:

[[2500]]

Classification Report:

	precision	recall	f1-score	support
cats	1.00	1.00	1.00	2500
accuracy			1.00	2500
macro avg	1.00	1.00	1.00	2500
weighted avg	1.00	1.00	1.00	2500

C:\Users\MGM\anaconda3\Lib\site-packages\sklearn\metrics\
_classification.py:409: UserWarning: A single label was found in 'y_true' and 'y_pred'.
For the confusion matrix to have the correct shape, use the 'labels' parameter to
pass all known labels. warnings.warn(